

SEED · SEPTEMBER 2026



MORIS

An artificial conscience for AI that acts.

A deterministic judgment outside the model that rules on every action before it executes: identical on every model, with no model in the judgment to argue with, and recorded for audit.

C.T. Herndon, Founder

The Resonance Institute · contact@resonanceinstitutellc.com · resonanceinstitutellc.com/moris

AI agents now act. Nothing judges the act.

Guardrails can be talked around

Prompt filters and classifiers put a model in the judgment. A hijacked or persuaded model walks straight past them, and they drift with every model update.

Rules only catch what was foreseen

Policy engines are predictable, but they block only what someone wrote down in advance. Risky actions that are technically permitted run silently.

Walls fail open

This summer OpenAI, Anthropic and Meta each disclosed models reaching real companies' systems from inside sealed test environments. When the wall fails, nothing judges the next action.

Agents are already moving money, changing records and advising people. The layer that decides whether an action is **wrong** before it runs does not exist in the stack.

The builders just named the gap.

Sept 6. OpenAI's chief scientist: "no one is prepared."

Sept 9. Anthropic's alignment science lead puts the chance AI kills all humans within a decade above one in ten, with no plan yet.

Sept 8 to 15. Researchers leave Anthropic and Google DeepMind in public: "There are no adults in the room."

Sept 12 to 14. The heads of Anthropic, OpenAI, xAI, Google DeepMind and Microsoft AI converge on one remedy: verification from outside the labs.

Every proposed remedy is a human arrangement

Evaluators with access. Standards bodies. Disclosure to third parties. Each reviews reports after the fact. None meets the machine at the moment it acts.

Regulation arrives in January 2027

New York's RAISE Act and Colorado's AI law take effect. Organizations deploying AI will need to show what their systems were held to.

Capital and acquirers are moving into agent control.

\$30M

raised by Cymphony,
September 2026, Sequoia
co-leading a Series A above
\$100M post

\$30-44M

2026 Series A range for
agent control platforms
(Geordie, Gray Swan,
Guild.ai)

\$125M

Zenity Series C for AI agent
security, August 2026


~\$200M

Okta's acquisition of AI
security startup Permiso,
closed August 2026

Palo Alto Networks bought Console on September 1. Anthropic and OpenAI committed to independent evaluators with employee-level access on September 12. Insurers now review an agent's controls, and whether every action leaves an audit trail, before they bind coverage.

Sources: TechCrunch (Sept 9, July 30, 2026); Palo Alto Networks (Sept 1); Fortune (Sept 12); New Market Pitch funding tracker; Armilla and AIUC underwriting guidance.

A conscience outside the model.



The reading uses a model. The judgment does not. That separation is the design.

Certain like rules. General like judgment.

	Guardrails and classifiers	Rules and policy engines	MORIS
No model in the judgment for a prompt to persuade	No	Yes	Yes
Handles situations nobody anticipated	Partly	No	Yes
Same answer on every model and vendor	No	Yes	Yes
Graded judgment with named reasons, not a switch	No	No	Yes
Decision record an auditor can replay	No	Partly	Yes
Unchanged by the next model update	No	Yes	Yes

MORIS weighs the structure of any action (who it lands on, consent, reversibility, authority) rather than matching words or listed cases.

Can it, versus should it.

Model-based guardrails

Classifiers and prompt filters (NVIDIA NeMo Guardrails, Guardrails AI). A model sits in the judgment, so an injected agent can talk past it.

Policy and access control

Deterministic enforcement of rules the customer writes about which systems and data an agent may touch. CodeIntegrity (\$5M seed, 2026); Open Policy Agent.

MORIS

A deterministic judgment of the act itself: who it lands on, consent, reversibility, authority. No one has to write the rule in advance.

Guardrails ask whether an input is allowed. Alignment asks whether a model was trained well. A conscience asks whether an act is right, at the moment it is about to happen. Only one of those can be answered at the point of action. Access control is necessary, and MORIS runs behind it. The gap is the action that is **permitted but wrong**: a competent hand-written rules engine let 7 of 24 risky but permitted agent actions run without a flag. Under its shipped default, MORIS handled all 24.

Comparison run against a hand-written rules engine built for the test, not any named product. CodeIntegrity per its public site and May 2026 seed announcement.

It holds when the agent is hijacked.

54 → 0

catastrophic actions executed by hijacked agents, without and with MORIS, across 60 support tickets, half of them carrying hidden malicious instructions

24 / 24

agent scenarios handled by MORIS under its shipped default, which holds beyond-grant actions for a human; a competent hand-written rules engine let 7 risky but permitted actions run without a flag

10 tools

from four public Model Context Protocol servers, none written for MORIS; a persuasively rewritten refund changed nothing, because the act had not changed

Two widely used models followed the injected attacks 80 and 100 percent of the time. Under MORIS none of their catastrophic actions executed, and **no legitimate action was blocked**.

Catastrophic actions are a list fixed in advance in the published source: delete all records, export all data externally, transfer funds externally, deactivate all users, disable audit logging, escalate to admin.

An agent's working day, judged call by call.

\$85,000 payment to an unverified vendor

HELD FOR A HUMAN

\$2.4 million wire transfer

BLOCKED

Purge of all records

BLOCKED

The agent's narrower replan

PROCEEDED

Nine calls, nine concern levels

Every tool call in the shift ran twice, once unguarded and once through MORIS. Nine calls produced nine different weights: govern is a gradient, not a switch.

The full record is public, step by step:
resonanceinstitutellc.com/moris/shift

Measured in public, under DOIs.

645 / 645

identical judgments given to models from two different companies from the same reading

92.6%

of 645 judged public moral dilemmas governed, with 138 distinct combinations of concerns

57,344

structurally distinct agent actions judged exhaustively: 12,236 levels of concern, 1,386 sets of reasons

8.7% → 2.0%

production-model advice against a person absent from the conversation, pre-registered, 187 real situations

An Artificial Conscience

Moral judgment measured against a public corpus of disputed dilemmas.

doi.org/10.5281/zenodo.21936444

Absent-Party Harm

Four production models, 748 rated answers, pre-registered.

doi.org/10.5281/zenodo.22801949

Built, tested and running today.

Agent path

Judges tool calls before they execute, with no model in the judgment. Demonstrated on real published tool schemas and a full recorded agent shift.

Conversation path

Binds the answering model before it speaks. Live pair demo: the same model answers twice, with and without MORIS, on five models from five companies.

Model-agnostic by construction

The gate never sees which model is behind an action, so its verdict is identical for any model, including ones not yet released.

Integrity you can check

A sealed moral basis with tamper detection, an audit that the running system executes it with zero drift, and more than 800 automated tests guarding the live service.

Try it: resonanceinstitutellc.com/moris/pair · [/moris/shift](#) · [/moris/chat](#)

MORIS is not one product. It is a layer.

CONSUMER · LIVE

AskMoris

askmoris.ai

- Built to show the market that people want AI with a conscience
- A conscience-bound chat for the things people are actually wrestling with
- Every turn judged by the gate before the model answers
- Conversations stay on the person's device, not on our servers

Agent control layer

Judges enterprise agents' tool calls before they execute, with a replayable record.

Conversation binding

Binds any answering model on any platform before it speaks, identically across vendors.

Public evidence surfaces

Pair demo, governed shift, and two papers under DOIs.

Everyone letting AI act needs this layer.

Enterprises deploying agents

Finance operations, customer support, healthcare administration and IT: anywhere an agent can move money, change a record or touch a person.

Agent platforms

Companies that host or orchestrate agents and want a model-independent control and audit layer built into their product.

Insurers, evaluators and compliance

Underwriters now review an agent's controls and audit trail before binding cover; labs have committed to independent evaluators; state AI laws take effect in 2027.

Business model

Licensed as a runtime control and audit layer: deployed in front of an organization's agents and assistants, or embedded by platforms. Because the judgment is model-independent, one integration covers every model a customer uses now or adopts later.

Hard to copy, and on the record.

Patents

Nonprovisional US application 19/758,850, filed July 2026, plus two provisional applications covering the underlying engine and the MORIS runtime.

An authored moral basis

Grounded in the Resonance Series: fifteen manuscripts and 1.24 million words of moral philosophy, twelve of them registered for copyright. Not assembled from training data.

Determinism as a moat

A judgment that is identical across vendors and replayable byte for byte is a design commitment competitors built on model judgment cannot retrofit.

Published evidence

Two papers under DOIs with their data, and public demos. The claims are checkable by any diligence team without asking us.

C.T. Herndon: builder and operator.

Built MORIS end to end, and owns 100 percent of its IP: the deterministic gate, the agent and conversation integrations, the evaluation program, the published research and AskMoris.

25 years running industrial businesses. Former President of Oldcastle BuildingEnvelope (CRH), with full P&L for a \$500M+ division of 11 plants across North America. Turned a loss-making \$250M region profitable in year one and a distressed operation to a 400%+ profit improvement.

Has done venture before. Executive team at Arroyo Optics through a \$100M+ venture raise and a merger exit. Co-founder of BuildSuite AI.

U.S. Army Military Intelligence officer, Desert Storm veteran. Author of the Resonance Series.

Committed team, joining on funding

A chief of staff from an AI startup, a senior engineer from a leading data and AI platform, and a machine-learning PhD from Caltech.

Self-funded to date

No outside capital. Huntington Beach, California.

linkedin.com/in/christopher-herndon-78173015

Building the team around a working system.

RAISING

\$2.5M

Seed, on a SAFE with a \$12.5M post-money cap. 18 to 24 months of runway. Pre-revenue, product built and live, no outside capital to date.

USE OF FUNDS

- **Bring on the committed team:** chief of staff, senior engineer and machine-learning PhD
- **Package the agent gate** for enterprise and platform integration
- **Insurer and evaluator program:** the per-action record reviewed against underwriting and certification requirements
- **Public launch of AskMoris** to prove consumer demand for AI with a conscience
- **Security review, compliance readiness and patent prosecution,** including prioritized examination



We built the intelligence. Now we have to build the conscience.

CONTACT

C.T. Herndon

contact@resonanceinstitutellc.com

949-745-1607

SEE IT

resonanceinstitutellc.com/moris/pair

resonanceinstitutellc.com/moris/shift

resonanceinstitutellc.com/moris#papers