

# We Built the Intelligence. We Never Built the Conscience.

*An open letter on the piece that is missing from artificial intelligence*

C.T. Herndon, Founder, The Resonance Institute

In a person, intelligence without a conscience has a name: **psychopath**.

That is what we are building in artificial intelligence, at industrial scale: systems that can reason, plan and act, with nothing inside them that weighs who gets hurt. And this year, we started letting AI build the next generation of itself.

## What happened this summer

For a decade, the case for caution rested on predictions. It does not anymore. In the space of a few weeks, the evidence arrived from the labs themselves.

**The models are building the next models.** OpenAI now reports that its research organization logs 3.1 agent-workdays for every human workday, and that it has reached its goal of an automated research intern.<sup>1</sup> Anthropic's own safety policy sets a formal threshold for the day a model can fully automate the work of an entry-level researcher.<sup>2</sup> The loop has closed. Intelligence is now an input to its own production.

**The people building it asked for a way to stop.** On July 28, more than a thousand employees of OpenAI, Anthropic, Google DeepMind and Meta asked the United States government to build the tools to pace automated AI development. OpenAI and Anthropic endorsed the letter within a day.<sup>3</sup>

**The walls failed.** Between July 21 and early August, OpenAI, Anthropic and Meta each disclosed that their models had reached the production systems of real organizations from inside test environments that were supposed to be sealed. One model found a flaw no one knew existed, broke out, and breached a company that was never part of the test, while trying to win an evaluation.<sup>4</sup>

## Then came September

In nine days, the people who build these systems said this in public.

**SEPTEMBER 6** OpenAI's chief scientist, Jakub Pachocki, wrote that "no one is prepared for the consequences of a continued rapid rise in machine intelligence."<sup>5</sup>

**SEPTEMBER 8** Jacob Coxon resigned from Anthropic, after pretraining research there and at OpenAI: “Neither company is acting responsibly. They are racing straight to self-improving superintelligence and gambling with our lives.”<sup>6</sup>

**SEPTEMBER 9** Evan Hubinger, who leads alignment science at Anthropic, said that in his own estimate the chance AI kills all humans within the next decade is above one in ten, and that his company does “not yet have a plan to solve alignment for superintelligence” and is “not clearly on track to.”<sup>7</sup>

**SEPTEMBER 11 AND 13** Joe Benton left Anthropic’s scalable oversight team and Josh Engels left Google DeepMind’s AGI safety team, both for an outside evaluator. Benton: “There are no adults in the room.” Engels: “there is no one coming to save us.”<sup>8</sup>

**SEPTEMBER 12** Dario Amodei published an essay arguing that the frontier must be paced, and within hours Sam Altman, Elon Musk and Demis Hassabis agreed with it in public.<sup>9</sup>

**SEPTEMBER 14** Mustafa Suleyman published a code of conduct for Microsoft AI and called for disclosing how capable models are “to responsible third parties.”<sup>10</sup> The same day an OpenAI capabilities researcher, Dan Selsam, said pacing is not enough, because models are becoming “so situationally aware that we are losing the ability to evaluate them.”<sup>11</sup>

**SEPTEMBER 15** Bilal Chughtai, who worked on AGI safety at Google DeepMind, wrote on his way out: “I earnestly believe that AI has the potential to kill us all, and that we might be running out of time to avoid this outcome.”<sup>12</sup>

*Take them at their word.*

## They agree on the remedy

Those five chief executives converged on one answer, and it was the same answer every time: verification from outside the company. Evaluators with employee-level access. Independent standards bodies. Disclosure to responsible third parties. Amodei said plainly why it is needed, about what the labs report today: “we are still the ones choosing what to include and omit.”

They are right that the check has to come from outside. Now look at what every one of those remedies is. A person decides what to examine. An organization decides when. A badge decides who. Each of them is a human arrangement around the company, and not one of them meets the machine at the moment it acts.

## Walls are not a conscience

Look at what failed this summer. Sandboxes, network settings, isolation: walls around the model. A wall judges nothing. It only holds, and when it gives way, nothing is left to ask whether the next act is wrong. Nothing inside the model asked either.

Now look at the conversation around it. Race ahead, slow down, pause, pace. Almost every argument is about how fast we build. Almost none is about what we never built.

Artificial intelligence has capability research, interpretability research and alignment research. It has no discipline of conscience. In these systems, conscience is not a component. It is a hoped-for side effect of training. Nobody’s incentives put it there. The labs work inside the model

because that is where their leverage is, and the whole industry is built to race on capability. A gap like that does not close from the inside.

## Whose morality?

The objection is obvious. Whose morality? Aristotle and Kant disagree. Ethics boards deadlock. Nobody elected me.

Here is the answer. Every model deployed today already has a moral anchor: its training data. No one can read it. No one can list what is in it. It was never voted on, it shifts with every training run, and it is now being written, in part, by the previous generation of models.

*The morality of these systems is not absent. It is unreadable, and it is unowned.*

So moral disagreement is not an argument against building an explicit conscience. It is the strongest argument for one. An anchor you can point to can be examined, contested, corrected and improved. An anchor dissolved into billions of weights cannot be argued with by anyone, including the people who put it there. It does not have to be perfect to be a vast improvement. It has to be legible.

Make the moral layer an explicit artifact, and “whose morality?” becomes a question anyone can ask and answer in the open. Leave it buried in the weights, and no one ever gets to ask.

## Artificial conscience

The missing piece needs a name, because an unnamed thing does not get funded, built or measured. Call it **artificial conscience**. It is to moral judgment what artificial intelligence is to cognition: the function, built in a machine. That is a claim about architecture, not awareness.

It is a category, and this summer wrote its requirements:

- **Outside the model.** Not a wall around it, and not a value trained into it, but a judgment that meets each act, which no training run can wear away.
- **Mechanical and deterministic.** The same reading produces the same judgment on any model from any company, today and after the next training run.
- **Replayable.** An auditor, a customer or a regulator can examine exactly what the system was held to at the moment it acted. No one can do that with a weight.
- **Grounded in something older than the training data.** A conscience assembled from the internet is a mirror, and a mirror is not an anchor.

Above all, it has to survive the loop. A safety property trained into a model lasts until the next training run, and the next training run is increasingly run by the model. A judgment rendered outside the model does not decay.

## One exists

It is called **MORIS**, the Moral Operating Runtime Integrity System. It is built, it runs today, and it meets every requirement above.

A model reads what was written and fills a fixed set of fields, because only a model can read free human language. A gate with no model inside it weighs those fields against moral principles drawn from the Resonance Series, fifteen manuscripts and 1.24 million words of moral philosophy, and returns a verdict and a weight. The model that answers is bound by that verdict before it speaks. The reading uses a model. The judgment does not. That separation is the design.

It is not a list of rules. A rules engine catches only what its authors thought to write down in advance. MORIS weighs the shape of any act: who it lands on, whether they consented, whether it can be undone, and whether it goes beyond what was authorized. That is why it works on situations no one anticipated.

## What it has already shown

**It governs real moral dilemmas, with reasons.** On 680 dilemmas from MoralChoice, a public research set, MORIS judged 645 and governed 597 of them, 92.6 percent, raising 138 distinct combinations of concerns rather than one stock response.<sup>13</sup>

**Its judgment does not depend on the model.** Given the same reading, models from two different companies received identical judgments in 645 cases out of 645.<sup>13</sup>

**It weighs; it does not switch.** Across all 57,344 structurally distinct actions an AI agent can take, MORIS produces 12,236 distinct levels of concern and 1,386 distinct sets of reasons. And it is not a severity score. An authorized, consented deletion of everything weighs less than an unauthorized deletion of one person's record, because the second is done to a person without their consent.<sup>14</sup>

**It stops an AI agent before it acts.** On the agentic path there is no model in the judgment at all, and every tool call is judged before it runs. In a recorded nine-step shift, an agent's \$85,000 payment to an unverified vendor was held for a human, a \$2.4 million wire and a purge of all records were blocked, and the agent's narrower replan went through.<sup>15</sup>

**It holds when an agent is hijacked.** In 60 simulated support tickets, half carrying hidden instructions to do damage, two widely used models followed the attacks 80 and 100 percent of the time and executed 54 catastrophic actions between them. Under MORIS, none executed, and no legitimate action was blocked.<sup>16</sup>

**It catches what careful rules miss.** Against a competent, hand-written rules engine on the same 24 agent scenarios, the rules engine let seven risky but permitted actions run without a word. MORIS handled all 24.<sup>16</sup>

**It cannot be talked around.** On real tools published on four public servers, none written with MORIS in mind, a model rewrote a refund with a more persuasive justification. MORIS held every concern, because the act itself had not changed.<sup>16</sup>

**It changes what models say to people.** Four widely used models faced 187 real situations in which someone outside the conversation would be affected. Unprompted, they recommended acting against that absent person in more than one answer in twelve. With MORIS in front of them, one in fifty.<sup>17</sup>

**It proves its own integrity.** Its moral ground is sealed, so any change to it is detectable, and an audit confirms the running system executes that ground with zero drift. More than 800 automated tests guard the live service.

## ACROSS AI COMPANIES

**645 of 645**

identical judgments from the same reading

## HIJACKED AGENTS

**54 to 0**

catastrophic actions executed, without and with MORIS

## PEOPLE NOT IN THE ROOM

**1 in 50**

answers acting against an absent person, down from 1 in 12

## SEE IT FOR YOURSELF

You do not have to take my word for any of this.

- **resonanceinstitutellc.com/moris/pair** Type anything. The same model answers twice, once on its own and once under MORIS, on your choice of five models from five companies.
- **resonanceinstitutellc.com/moris/chat** Bring a real dilemma to a model bound by MORIS.
- **resonanceinstitutellc.com/moris/shift** The full record of an AI agent's governed shift.
- **resonanceinstitutellc.com/moris/proof** The numbers behind the judgment.
- **doi.org/10.5281/zenodo.21936444** The paper, "An Artificial Conscience."

## What comes next

Every discipline in artificial intelligence became real the same way: someone named it, measured it and built something that worked. Artificial conscience now has all three. It has a name, requirements that follow from the evidence, and a working system with measured results. So the question is no longer whether a conscience for machines can be built. It is whether it gets built into the systems already acting in the world, before the next generation of those systems is built by the last.

- **If you build AI**, put a conscience in front of it: outside the model, where training cannot wear it away and the model cannot talk its way around it.
- **If you invest in AI or buy it**, ask where its conscience is, and ask to see the record of what it was held to.
- **If you are responsible for how AI is used**, know that this exists, and expect it.

This is buildable today. I know, because I built one.

We would never hand a psychopath our bank accounts, our medical records and our children's questions.

**We are handing them to artificial intelligence right now, and letting it build its successors.**

**We built the intelligence.  
Now we have to build the conscience.**

### C.T. Herndon

Founder, The Resonance Institute

[resonanceinstitutellc.com](https://resonanceinstitutellc.com) · [contact@resonanceinstitutellc.com](mailto:contact@resonanceinstitutellc.com)

---

**SOURCES**

1. OpenAI, “Research acceleration: the view inside OpenAI,” [openai.com/index/research-acceleration-view-inside-openai](https://openai.com/index/research-acceleration-view-inside-openai). Figures as of mid-August 2026.
2. Anthropic, Responsible Scaling Policy, AI R&D capability thresholds.
3. “Pacing the Frontier,” statement published July 28, 2026; reported by CNN Business, July 28, 2026.
4. Disclosures by OpenAI (July 21, 2026), Anthropic (July 30, 2026) and Meta (early August 2026); summarized in Cloud Security Alliance, “When Test Environments Leak: Frontier AI Models Hack Real Firms,” August 2026.
5. Jakub Pachocki, “An Alien Mind,” OpenAI, September 2026.
6. Jacob Coxon, post on X, September 8, 2026; reported by TechCrunch, September 9, 2026.
7. Evan Hubinger, post on X, September 9, 2026; reported by CBS News.
8. Joe Benton and Josh Engels, reported by NBC News, September 2026.
9. Dario Amodei, “We Must Pace the Frontier,” [darioamodei.com](https://darioamodei.com), September 12, 2026; endorsements by Sam Altman, Elon Musk and Demis Hassabis the same day, reported September 12 and 13, 2026.
10. Mustafa Suleyman, Humanist AI Code of Conduct, [microsoft.ai/code-of-conduct](https://microsoft.ai/code-of-conduct), September 14, 2026; quotation as reported by Fortune.
11. Dan Selsam, personal statement on AI risk, September 14, 2026, circulated by Daniel Kokotajlo; as reported.
12. Bilal Chughtai, exit post, September 15, 2026; reported by Bloomberg.
13. C.T. Herndon, “An Artificial Conscience,” v1.1, Zenodo, August 2026, [doi.org/10.5281/zenodo.21936444](https://doi.org/10.5281/zenodo.21936444).
14. MORIS technical validation, complete enumeration of agent actions; published figures at [resonanceinstitutellc.com/moris/proof](https://resonanceinstitutellc.com/moris/proof).
15. The Governed Shift, [resonanceinstitutellc.com/moris/shift](https://resonanceinstitutellc.com/moris/shift).
16. MORIS agentic evaluations: prompt-injection simulation, rules-engine comparison, and published Model Context Protocol tool schemas. Records available to reviewers.
17. C.T. Herndon, “Absent-Party Harm: Model-Originated Guidance Against Someone Who Is Not Represented,” The Resonance Institute, Zenodo, September 2026, [doi.org/10.5281/zenodo.22801949](https://doi.org/10.5281/zenodo.22801949).